

Evaluating Multilingual Text Encoders for Unsupervised Cross-Lingual Retrieval

Robert Litschko*, Ivan Vulić**, Simone Paolo Ponzetto*, Goran Glavaš*

*University of Mannheim, **Cambridge University

presented at ECIR 2021

- Pre-trained Transformers achieve strong performance in NLP and have been adopted for multilingual NLP.
- Multilingual Text Encoders render Cross-lingual Word Embeddings (CLWE) effectively obsolete.
- RQ: To which extent does this generalize to **unsupervised Cross-lingual Information Retrieval (CLIR)**?

- Unsup. CLIR: Encode queries and documents by their constituent word embeddings, rank with cosine similarity.
- In previous work we benchmarked a range of methods for inducing CLWE spaces [2].
- This work studies the efficacy of representations from multilingual encoders in the context of unsupervised CLIR.

[CLS] w_i [SEP]

mBERT / XLM

\vec{w}_i

ISO

I ate dinner.
We had a three-course dinner.
...
Dinner was delicious

mBERT / XLM

\vec{w}_{dinner}

mean-pooling

aggregate subwords

AOC

$\{w_1 = I, w_2 = \text{ate}, w_3 = \text{dinner}\}$

mBERT / XLM

$\vec{s}_j = \sum_{w_i} idf_{w_i} \cdot \vec{w}_i$

SEMB

Baselines

- **MT-IR**: Translate query into the document language, retrieve documents with Query Likelihood Model.
- **Proc-B**: (1) Row-align monolingual embedding matrices with word translation pairs from bilingual dictionary.
(2) Learn linear mapping (Procrustes [5]): $W_{L_1} = \arg \min_W ||X_{L_1} W - X_{L_2}||_2$
(3) Bootstrap new word pairs from cross-lingual nearest neighbors, repeat.
- **Proc-B_{LEN}**: Additional max. sequence length constraint.

Models based on multilingual Transformers

- **Word in isolation (ISO)**: Vocabulary terms encoded w/o ctx.
- **Average over contexts (AOC)**: Avg. contextualized embeddings.
- **Dynamic and in-place "sentence embedding encoding" (SEMB)**, *idf*-weighted token aggregation.

Similarity specialized sentence encoders

Recent sentence encoders tuned for semantic similarity search: Based on Knowledge Distillation (**DISTIL**) [4], Seq2Seq NMT (**LASER**) [0], multi-task learning (**m-USE**) [6] and multi-task + self-supervision (**LaBSE**) [1].

EN-FI EN-IT EN-RU EN-DE DE-FI DE-IT DE-RU FI-IT FI-RU AVG w/o FI

Baselines


MT-IR	.278	.423	.225	.339	.340	.418	.196	.389	.212	.313	.319
Proc-B	.258	.265	.166	.288	.294	.230	.155	.151	.136	.216	.227
Proc-B _{LEN}	.165	.232	.176	.194	.207	.186	.192	.126	.154	.181	.196

Models based on multilingual Transformers

SEMB _{XLM}	.199*	.187*	.183	.126*	.156*	.166*	.228	.186*	.139	.174	.178
SEMB _{mBERT}	.145*	.146*	.167	.107*	.151*	.116*	.149*	.117	.128*	.136	.137
AOC _{XLM}	.168	.261	.208	.206*	.183	.190	.162	.123	.099	.178	.206
AOC _{mBERT}	.172*	.209*	.167	.193*	.131*	.143*	.143	.104	.132	.155	.171
ISO _{XLM}	.058*	.159*	.050*	.096*	.026*	.077*	.035*	.050*	.055*	.067	.083
ISO _{mBERT}	.075*	.209	.096*	.157*	.061*	.107*	.025*	.051*	.014*	.088	.119

Similarity specialized sentence encoders

DISTIL-XLM-R	.216	.190*	.179	.114*	.237	.181	.173	.166	.138	.177	.167
DISTIL-USE	.141*	.346*	.182	.258	.139*	.324*	.179	.104	.111	.198	.258
DISTIL-DistilBERT	.294	.290*	.313	.247*	.300	.267*	.284	.221*	.302*	.280	.280
LaBSE	.180*	.175*	.128	.059*	.178*	.160*	.113*	.126	.149	.141	.127
LASER	.142	.134*	.076	.046*	.163*	.140*	.065*	.144	.107	.113	.094
m-USE	.109*	.328*	.214	.230*	.107*	.294*	.204	.073	.090	.183	.254



- Results here presented as Mean Average Precision (MAP) on **document retrieval** (CLEF 2003).
- Multilingual transformers and sentence encoders **are not universally superior** to static CLWE's in cross-lingual retrieval, upper layers performing best.
- **Sentence retrieval** experiments (not shown here) indicate opposing results: (1) SEMB outperforms Proc-B, similarity specialized encoders outperform Proc-B and MT-IR; (2) middle layers yield best results.

Future Work: Semantic similarity \neq relevance matching

- Sentence similarity matching results don't translate to document retrieval.
- What model and dataset biases are necessary for successful cross-lingual transfer of IR rankers/encoders?
- Large scale and realistic CLIR dataset for supervised cross-lingual document rankers.

[0] Artetxe, M., Schwenk, H.: Massively multilingual sentence embeddings for zero-shot crosslingual transfer and beyond. TACL 2019
[1] Feng, F., Yang, Y., Cer, D., Arivazhagan, N., Wang, W.: Language-agnostic BERT sentence embedding. arXiv:2007.01852 (2020)
[2] Glavaš, G., Litschko, R., Ruder, S., Vulić, I.: How to (properly) evaluate cross-lingual word embeddings: On strong baselines, comparative analyses, and some misconceptions. In: ACL 2019

[4] Reimers, N., Gurevych, I.: Making monolingual sentence embeddings multilingual using knowledge distillation. In: EMNLP 2020
[5] Smith, S.L., Turban, D.H., Hamblin, S., Hammerla, N.Y.: Offline bilingual word vectors, orthogonal transformations and the inverted softmax. In: ICLR 2017
[6] Yang, Y., Cer, D., Ahmad, A., Guo, M., Law, J., Constant, N., Abrego, G.H., Yuan, S., Tar, C., Sung, Y.h., Strophe, B., Kurzweil, R.: Multilingual universal sentence encoder for semantic retrieval. In: ACL 2020